



MongoDB and Cloud Storage

Frank Walsh, Diarmuid O'Connor

Agenda

- Cloud Databases
- MongoDB
 - Querying
 - Integrating with Node.js
 - The Contacts API implementation

Databases in Enterprise Apps

- Most data driven enterprise applications need a database
- In traditional enterprise applications, this requires
 - Backups
 - Fail over
 - Maintenance
 - Capacity provisioning
- Usually handled by a Database Administrator.

Databases in the Cloud

- For some apps, a traditional database may not be the best fit
 - Does the app require transactional integrity
 - Do you need db schema definition
 - Do you know your scaling requirements, particularly if it's a web app
- One approach is to use the **Cloud** for you DB
 - Designed for scale
 - Can be outsourced so you don't have to deal with infrastructure requirements.

Cloud DB Advantages

- Removes Management costs
- Inherently scalable
- Latency is predictable and constant
- No need to define schemas etc.
- Lots of Cloud DB offerings out there
 - SQL based
 - NoSQL based
- If organisation policy/standards do not allow outsourcing:
 - Can host yourself, most NoSQL DBs are free.

Cloud Database Practices

- Drop Consistency
 - this makes distributed systems much easier to build
- Drop SQL and the relational model
 - simpler structures are easier to distribute:
 - key/value pairs
 - **structured documents**
 - **pseudo-tables**
 - tend to be schema-free, accepting data as-is
- Offer HTTP interfaces using XML or **JSON**
- Use in-memory storage aggressively

Designing Distributed Data

- App data is not homogeneous
 - some kinds of data will be much larger
- consider using different databases for different requirements:
- user details,billing - needs consistency
 - require traditional database
- user data,content - needs partition tolerance
 - replicate to keep safe
- analytics,sessions - needs availability
 - "eventually consistent" is good enough

MONGODB

Introduction

- Document-oriented database
 - but closer to traditional SQL databases than others
- Uses JSON natively - perfect fit for Node.js
- Query language with many SQL features
 - Uses JSON too, and has an easy learning curve
- Commercial support - 10gen.com product
 - cloud hosting providers - e.g. mongoLab.com
- Community support - popular choice

Mongo Terminology

- Each database contains a set of "Collections"
- Collections are analogous to SQL tables
- Collections contain a set of JSON documents
 - there is no schema (in the DB)
- the documents can all be different
 - means you have rapid development
 - adding a property is easy - just starting using in your code
- makes deployment easier and faster
 - roll-back and roll-forward are safe - unused properties are just ignored
- Collections can be indexed and queries
- Operations on individual documents are atomic

Getting Started (locally)

- For complete MongoDB installation instructions, see [the manual](#).
- Starting MongoDB:
`mongod`
- This starts the process.
- Can add other parameters, for instance location of data.

Mongo Shell

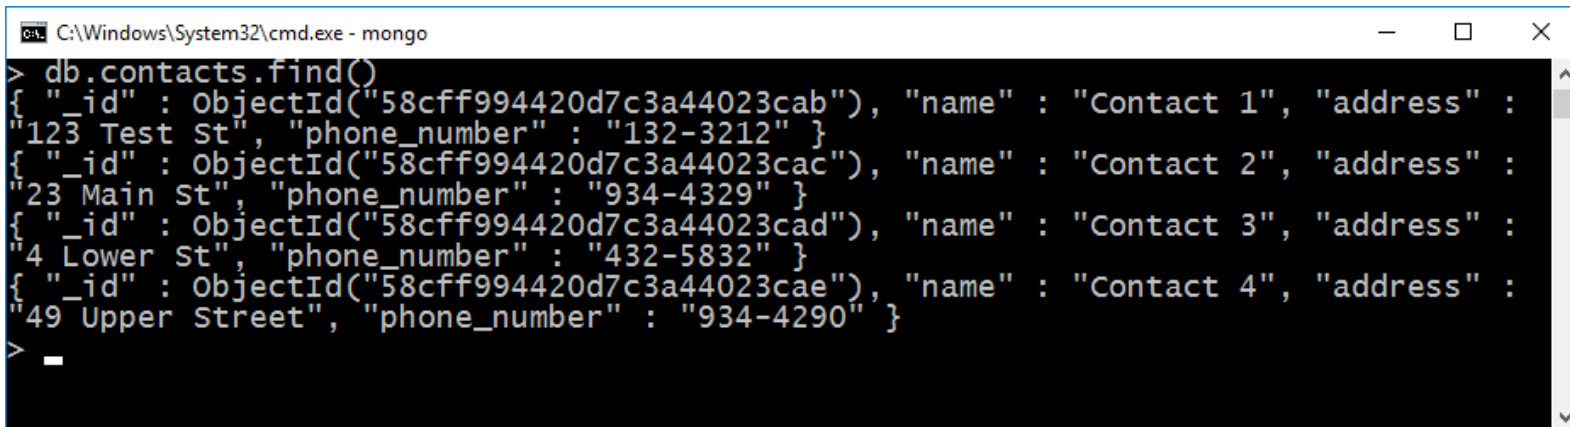
- Interactive JavaScript interface to MongoDB.
- Query/update data and perform administrative operations.

```
C:\repos\webervicesdev-2017>mongo
MongoDB shell version v3.4.2
connecting to: mongodb://127.0.0.1:27017
MongoDB server version: 3.4.2
Server has startup warnings:
2017-03-20T13:49:38.768+0000 I CONTROL [initandlisten]
2017-03-20T13:49:38.769+0000 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2017-03-20T13:49:38.773+0000 I CONTROL [initandlisten] **           Read and write access to data and configuration is u
nrestricted.
2017-03-20T13:49:38.775+0000 I CONTROL [initandlisten]
>
```

- By default, Mongo shell will attempt to connect to the MongoDB instance running on the localhost interface on port 27017.

The MongoDB Query Language

- MongoDB provides a JavaScript API and JSON-based query language
- Use the MongoDB shell to execute queries
 - similar to using MySQL console
- Example: list of contacts

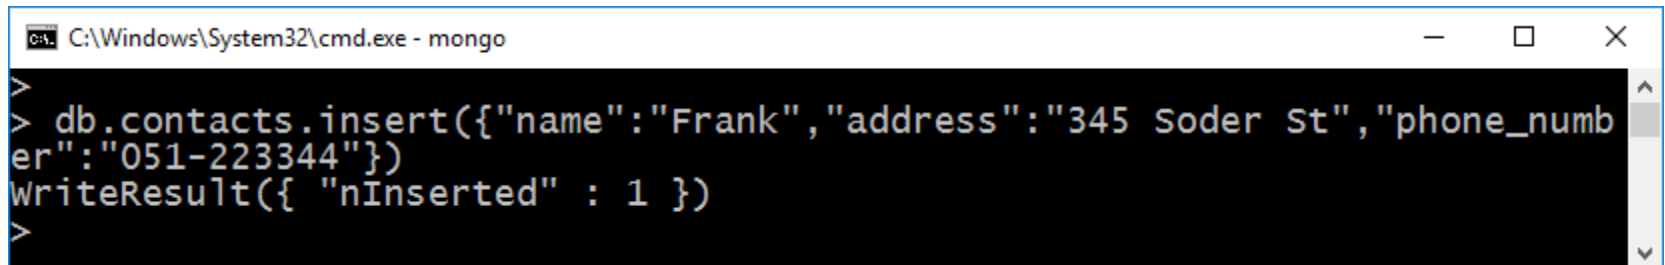


```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.find()
{ "_id" : ObjectId("58cff994420d7c3a44023cab"), "name" : "Contact 1", "address" :
123 Test St", "phone_number" : "132-3212" }
{ "_id" : ObjectId("58cff994420d7c3a44023cac"), "name" : "Contact 2", "address" :
23 Main St", "phone_number" : "934-4329" }
{ "_id" : ObjectId("58cff994420d7c3a44023cad"), "name" : "Contact 3", "address" :
4 Lower St", "phone_number" : "432-5832" }
{ "_id" : ObjectId("58cff994420d7c3a44023cae"), "name" : "Contact 4", "address" :
49 Upper Street", "phone_number" : "934-4290" }
>
```

- db - current database
- contacts - the contacts collection
- .find() - collection API method (corresponds to collection URL in last lecture...)
- The Result Set is a list of JavaScript objects, representing matched documents

MongoDB: Inserts

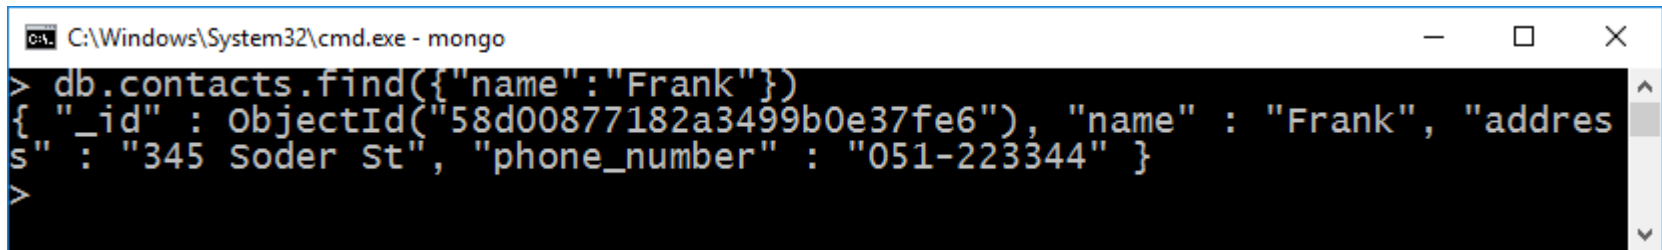
- Collections do not need to be created explicitly
 - just insert a document
- MongoDB automatically assigns a 12 byte unique identifier to any document
 - the **_id** property
 - Stored internally as binary

A screenshot of a Windows command prompt window titled "C:\Windows\System32\cmd.exe - mongo". The window has standard Windows window controls (minimize, maximize, close) in the top right corner. The command prompt shows a MongoDB shell session with the following text:

```
> db.contacts.insert({"name":"Frank","address":"345 Soder St","phone_number":"051-223344"})
WriteResult({ "nInserted" : 1 })
>
```

MongoDB:Queries

- Documents are retrieved by specifying a set of conditions to match against
- simplest case : query-by-example
- provide a subset of properties that must match

A screenshot of a Windows command prompt window titled "C:\Windows\System32\cmd.exe - mongo". The prompt shows a MongoDB query: > db.contacts.find({"name":"Frank"}) followed by a JSON document: {"_id": ObjectId("58d00877182a3499b0e37fe6"), "name": "Frank", "address": "345 Soder St", "phone_number": "051-223344" }.

```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.find({"name":"Frank"})
{"_id" : ObjectId("58d00877182a3499b0e37fe6"), "name" : "Frank", "address" : "345 Soder St", "phone_number" : "051-223344" }
```

- More complex queries use a convention of embedded meta-properties to specify conditions these are signified with a \$ prefix.
 - Example:{name:{\$exists:true}}
 - returns documents that have a name property

MongoDB: Queries

- Common meta-properties used to query data are:
 - \$gt, \$gte, \$lt, \$lte** meaning $>$, \geq , $<$, \leq

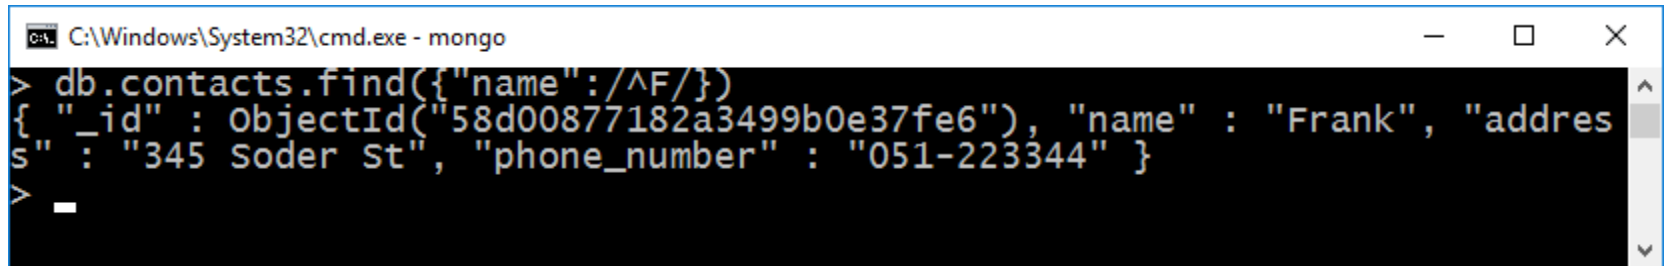
```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.insert({"name":"Mary","age":21,"address":"345 Soder St","phone_number":"051-223344"})
WriteResult({ "nInserted" : 1 })
> db.contacts.insert({"name":"Jane","age":31,"address":"345 Keel St","phone_number":"051-445566"})
WriteResult({ "nInserted" : 1 })
> db.contacts.find({"age":{"$gte":21,$lt:31}})
{"_id" : ObjectId("58d00909182a3499b0e37fe7"), "name" : "Mary", "age" : 21, "address" : "345 Soder St", "phone_number" : "051-223344" }
```

- \$or, \$in, \$nin**

```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.find({"name":{"$in":["Jane","Frank"]}})
{"_id" : ObjectId("58d00877182a3499b0e37fe6"), "name" : "Frank", "address" : "345 Soder St", "phone_number" : "051-223344" }
{"_id" : ObjectId("58d00939182a3499b0e37fe8"), "name" : "Jane", "age" : 31, "address" : "345 Keel St", "phone_number" : "051-445566" }
```


MongoDB: Queries

- **regular expressions** {word: /th^/i }

A screenshot of a Windows command prompt window titled "C:\Windows\System32\cmd.exe - mongo". The prompt shows a MongoDB query: > db.contacts.find({"name":/^F/}) followed by a JSON document: {"_id": ObjectId("58d00877182a3499b0e37fe6"), "name": "Frank", "address": "345 Soder St", "phone_number": "051-223344" }. The prompt is followed by a cursor and a hyphen.

```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.find({"name":/^F/})
{"_id": ObjectId("58d00877182a3499b0e37fe6"), "name": "Frank", "address": "345 Soder St", "phone_number": "051-223344" }
> -
```

- db.contacts.find().limit(5)
 - limits the number of documents in the result set.
- db.contacts.find().skip(5)
 - Set the Starting Point of the Result Set

MongoDB:Updates

- Documents are updated by providing:
 - a query to select the relevant subset of documents,
 - an update specification, which is either:
 - a complete replacement document
 - meta-properties that modify specific document properties
- example:
\$set changes specific properties
Example:complete replacement:
> db.city.insert({name:'dublin'})
> db.city.update({name:'dublin'}, {name:'Dublin',county:'Dublin'})
- Example:modify specific properties:
> db.city.insert({name:'Cork',county:'cork'})
> db.city.update({name:'Cork'}, {\$set:{county:'Cork'}})
- See <http://www.mongodb.org/display/DOCS/Updating> for more

MongoDB:Update Properties

- Common meta-properties used with the update command are:
 - **\$set** - sets specified properties, but leaves others alone
\$set:{name:'New Name'}
- **\$unset** - deletes specified properties
\$unset:{name:1}
- **\$inc** - increments a numeric property
inc:{ upvotes: 2 }
adds 2 to the counter property, or if it does not exist, sets it to 2
- **\$push, \$pop** - add to or remove values from, an array
 - \$push: { comments: {who:..., msg:...} }
 - \$pop: {comments: -1 }

MongoDB:Upserts

- The MongoDB update command can optionally insert a document if it is not found. This is known as an 'upsert'
- This is useful when starting counters as it avoids corrupting the count when two independent updates try to initialize the counter.

```
db.counters.update( {name:'foo'}, {$inc:{value:1}},  
true)
```

- The first update will create the counter:
{name:'foo', value:1}
- The second update will increment the counter:
{name:'foo', value:2}

MONGOOSE

Mongoose Overview

- Mongoose is a object-document model module in Node.js for MongoDB
 - Wraps the functionality of the native MongoDB driver
 - Exposes models to control the records in a doc
 - Supports validation on save
 - Extends the native queries

Installing Mongoose

- Run the following from the CMD/Terminal
 - \$ npm install -save mongoose
- In node
 - Load the module
 - import mongoose from 'mongoose';
- Connect to the database
 - mongoose.connect(mongoDbPath);

Mongoose Schemas and Models

- Mongoose supports models
 - i.e. fixed types of documents
 - Needs a mongoose.Schema
 - Each of the properties must have a type
 - Number, String, Boolean, array, object

```
1  const mongoose = require('mongoose'),
2  Schema = mongoose.Schema;
3
4  ▼ const ContactSchema = new Schema({
5      name: String,
6      address: String,
7      age: Number,
8      email: String,
9      updated: Date
10  });
11
12  const ContactModel = mongoose.model('contacts', ContactSchema);
```


Mongoose Schemas - Arrays

Comments property is
an Array of
CommentSchemas

```
1 |const mongoose = require('mongoose'),
2 |Schema = mongoose.Schema;
3
4 |const CommentSchema = new Schema({
5 |  body: {type: String, required:true},
6 |  author: {type: String, required:true},
7 |  upvotes:Number
8 |});
9
10|const PostSchema = new Schema({
11|  title: {type: String, required:true},
12|  link: {type: String, optional:true},
13|  username: {type: String, required:true},
14|  comments: [CommentSchema],
15|  upvotes: { type: Number, min: 0, max: 100 }
16|});
17
18|export default mongoose.model('posts', PostSchema);
```

Mongoose Schema - Validation

- Can define validation constraints on properties :

```
1  const mongoose = require('mongoose'),
2  Schema = mongoose.Schema;
3
4  ▼ const ContactSchema = new Schema({
5    name: String,
6    address: String,
7    age: { type: Number, min: 0, max: 120 },
8    email: String,
9    updated: { type: Date, default: Date.now }
10  });
11
12  const ContactModel = mongoose.model('contacts', ContactSchema);
```

Mongoose Custom Validation

- Developers can define custom validation on their properties (e.g. validate length of comment when trying to save)

```
18 CommentSchema.path('body').validate((v)=>{
19     if (v.length>40 || v.length < 5){
20         return false
21     }
22     return true
23 })
24
```

Data Manipulation Mongoose

- Mongoose supports all the CRUD operations:
 - Create → `Model.create()`
 - Read → `Model.find()`
 - Update → `Model.update(condition, props, cb)`
 - Remove → `Model.remove()`
- Can operate with "*error first*" callback or promises.

Create Contact with Mongoose

```
import mongoose from 'mongoose';
const Schema = mongoose.Schema;

const ContactSchema = new Schema({
  name: String,
  address: String,
  age: {
    type: Number,
    min: 0,
    max: 120,
  },
  email: String,
  updated: {
    type: Date,
    default: Date.now,
  },
});

export default mongoose.model('Contact', ContactSchema);
```

```
// Create a contact, using async handler
router.post('/', asyncHandler(async (req, res) => {
  const contact = await Contact.create(req.body);
  res.status(201).json(contact);
}));
```

Update Contact with Mongoose

```
// Update a contact
router.put('/:id', asyncHandler(async (req, res) => {
  if (req.body._id) delete req.body._id;
  const contact = await Contact.update({
    _id: req.params.id,
  }, req.body, {
    upsert: false,
  });
  if (!contact) return res.sendStatus(404);
  return res.json(200, contact);
}));
```

Mongoose Queries

- Mongoose provides a more expressive version of the native MongoDB
 - Instead of:
`{ $or: [{ conditionOne: true }, { conditionTwo: true }] }`
 - Do:
`.where({ conditionOne: true }).or({ conditionTwo: true })`

Mongoose Queries

- Mongoose supports many queries:
 - For equality/non-equality
 - Selection of some properties
 - Sorting
 - Limit & skip
- All queries are executed over the object returned by `Model.find*()`
 - `Model.findOne()` returns a single document, the first match
 - `Model.find()` returns all
 - `Model.findById()` queries on the `_id` field.

```
// Delete a contact
router.delete('/:id', asyncHandler(async (req, res) => {
  const contact = await Contact.findById(req.params.id);
  if (!contact) return res.send(404);
  await contact.remove();
  return res.status(204).send(contact);
}));
```


Mongoose Queries

- Can build complex queries and execute them later

```
1  const query = ContactModel.where('age').gt(17).lt(66)
2    .where('county').in(['Waterford', 'Wexford', 'Kilkenny']);
3
4  query.exec((err, contacts) => {...})
5
6
```

- The above finds all contacts where age >17 and <66 and living in either Waterford, Kilkenny or Wexford

Mongoose Sub-Docs

- Ex: Hacker News – Adding a comment to a post.

```
// add comment
router.post('/:id/comments', asyncHandler( async (req, res) => {
  const id = req.params.id;
  const comment = req.body;
  const post = await Post.findById(id);
  post.comments.push(comment);
  await post.save();
  return res.status(201).send({post});
}));
```

Mongoose Sub-Docs

- Updating a Sub-Document(e.g. incrementing the upvotes for a comment)

```
router.post('/:postId/comments/:commentId/upvotes', asyncHandler( async (req, res) => {  
  const commentId = req.params.commentId;  
  const postId = req.params.postId;  
  const post = await Post.findById(postId);  
  post.comments.id(commentId).upvotes++;  
  await post.save();  
  return res.status(201).send({post});  
}));
```

Each subdocument is assigned
it's own `_id` from MongoDB.
This is a special method to
access sub documents

Mongo Sub docs

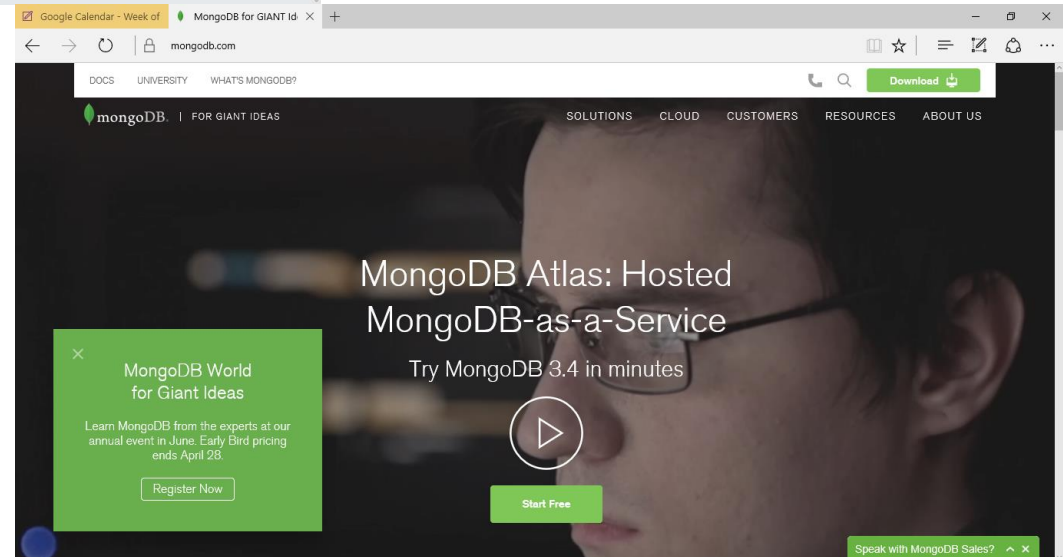
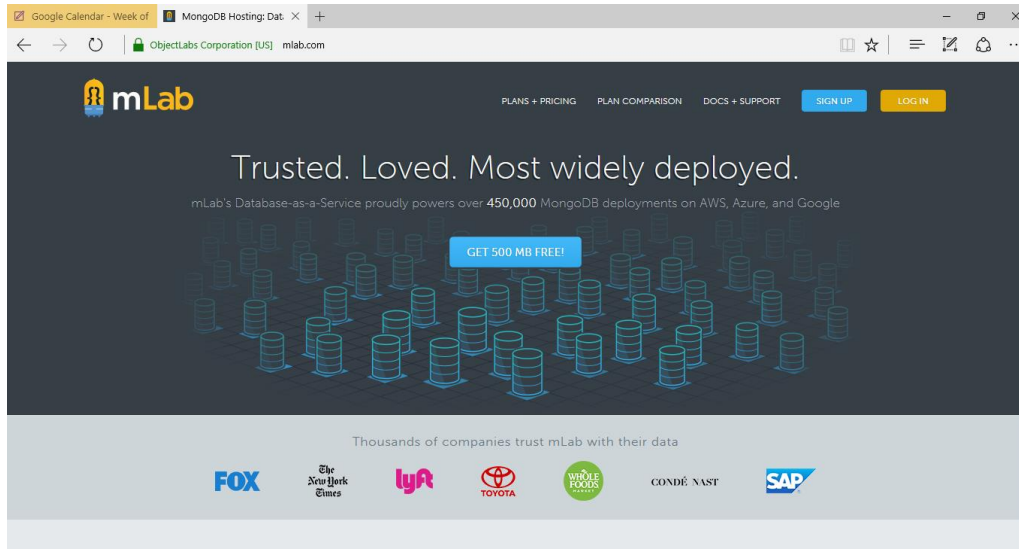
- Removing a sub document

```
router.delete('/:postId/comments/:commentId', asyncHandler( async (req, res) => {  
  const commentId = req.params.commentId;  
  const postId = req.params.postId;  
  const post = await Post.findById(postId);  
  post.comments.id(commentId).remove();  
  await post.save();  
  return res.status(201).send({post});  
}));
```

MongoDB as a Service


- Best practice for initial development is to host MongoDB process on your development machine
- In production environments, Mongo will be hosted:
 - on it's own instance or
 - provisioned as a service

MongoDB as a Service



MongoDB as a Service

- Most providers allow free access tier
- Provide user credentials wrapped in a URL
- All you need to do is update your config with the relevant URL
- Careful to ignore credentials when pushing to github/public repo



[Home](#)
Database: contacts_db

To connect using the mongo shell:

```
% mongo ds039311.mlab.com:39311/contacts_db -u <dbuser> -p <dbpassword>
```

To connect using a driver via the standard MongoDB URI ([what's this?](#)):

```
|mongodb://<dbuser>:<dbpassword>@ds039311.mlab.com:39311/contacts_db
```